

Working title: InterviewOS

The Mock Interview Platform Built on Real Interviews

A voice-first AI mock interview platform grounded in real, role-specific and company-specific interview patterns — covering product companies, service-based IT firms, European employers, and non-engineering corporate functions with equal seriousness.

Document	Product Requirements Document (PRD)
Version	v1.1 — session-scoped interview model
Date	25 August 2026
Owner	Naman Kaushik
Phase 1 platform	Responsive web application
Phase 3 platform	Native mobile (Android first, iOS second)
Status	Draft for build planning

The one-line thesis

Every AI interview tool on the market prepares you for a product-company system design round. We do that too, and we intend to do it well. But we also prepare you for a Deloitte manager round, an Infosys client-fit discussion, a TCS techno-managerial panel, a European culture-fit loop, a CA firm articleship interview, and a QA automation screen — conducted in the language and register those interviews actually happen in. Breadth without being worse at the narrow case everyone else already serves. That is the product.

Contents

01	Problem statement and market gap
02	Product thesis and defensible moat
03	Target segments and coverage map
04	Data strategy: what we source and how
05	The candidate profile model
06	Interview engine architecture
07	Voice pipeline: design and unit economics
08	RAG and knowledge layer
09	Feedback and scoring system
10	Website: information architecture and design
11	Technology stack
12	Phased delivery roadmap
13	Web-to-mobile transition strategy
14	Monetisation and unit economics
15	Success metrics
16	Risks, legal posture and mitigations
17	Open decisions

Problem statement and market gap

Interview preparation tooling has converged on a narrow slice of the market. The current generation of AI mock-interview products optimises almost exclusively for data structures, algorithms and distributed system design, aimed at candidates targeting large product companies in the US market. This is a well-served niche and a crowded one.

The under-served reality is considerably larger:

- **Service-based IT firms** (Deloitte, Infosys, TCS, Accenture, Cognizant, Wipro, Capgemini, LTIMindtree) hire at enormous volume, and their loops look nothing like a product-company loop. They weight client-scenario handling, project walkthroughs, process knowledge, techno-managerial rounds and HR/values fit far more heavily than algorithmic puzzles.
- **European employers** run culture-fit, structured competency and values-based rounds that Indian candidates are systematically under-prepared for, alongside visa/relocation conversations that no existing tool rehearses.
- **Non-engineering corporate functions** — product management, QA, data and AI engineering, chartered accountancy, taxation, management consulting, supply chain, automobile and manufacturing, HR, sales — have effectively zero credible AI mock interview coverage, despite representing the majority of white-collar hiring.
- **Language and register.** A very large share of interviews in the Indian market are conducted in code-switched Hindi-English. Every existing tool conducts a sanitised American-English interview, which is a poor rehearsal for the real thing.

Where existing tools break down

They generate questions from a job title. That produces plausible-sounding but generic questions, divorced from how a specific employer actually structures its rounds, what it probes on, and what a real interviewer does when a candidate gives a weak answer. Candidates report the experience feels like a quiz, not an interview.

Product thesis and defensible moat

The platform is a voice-first, adaptive mock interviewer that conducts a realistic, conversational interview — personalised to the candidate's resume, their target employer, the specific round, and the language they will actually be interviewed in — and returns a structured, evidence-backed feedback report.

Three layers of defensibility

A content moat built on scraped question banks is not achievable for a new entrant. GeeksforGeeks, LeetCode and Blind have a decade-plus head start on community volume, and attempting to out-scrape them carries legal exposure without producing differentiation. The moat is therefore deliberately placed elsewhere.

Layer	What it is	Why a competitor cannot copy it quickly
1. Interview realism	Adaptive voice interviewer with a consistent persona, natural follow-ups, interruption handling, pressure and silence — not a question list read aloud.	Requires sustained tuning of conversation design, latency engineering and persona work. This is craft, not a dataset, and it compounds with every session.

Layer	What it is	Why a competitor cannot copy it quickly
2. Coverage breadth	Service-based firms, European employers and non-engineering functions, at genuine parity with product-company coverage — which remains first-class, not an afterthought.	Incumbents are structurally anchored to the DSA/system-design market and their brand, content and funnel all point there. Broadening is expensive for them.
3. Language and cultural fit	Hindi-English code-switched interviews as a first-class mode, plus register and etiquette matched to the employer type.	Requires native fluency in the interview culture being simulated. Foreign competitors cannot fake this convincingly and generally do not attempt it.

Sequencing note on the moat

Layers 1 to 3 are available from day one and require no user base. The community-contributed interview corpus (Section 04) is a fourth layer that compounds over years — it is the long-term defensibility, but it must not be a launch dependency. The product has to be excellent before a single user contributes data.

Target segments and coverage map

Coverage is modelled along two axes: **employer archetype** (which determines round structure and evaluation criteria) and **function** (which determines subject matter). At the start of each session the candidate names the company and role they want to be interviewed for; the platform resolves that company to an archetype and composes the appropriate loop. Selection is per session rather than a stored target list — see Section 05.

Employer archetypes

Archetype	Representative employers	Round structure emphasis
Global product	Google, Amazon, Microsoft, Atlassian, Uber	Coding, system design, behavioural/leadership-principle rounds, bar raiser
Indian product / startup	Zoho, Freshworks, Razorpay, Zerodha, Swiggy	Practical coding, product sense, ownership scenarios, founder round
Service-based IT	TCS, Infosys, Wipro, Cognizant, Capgemini, LTIMindtree	Aptitude, technical fundamentals, project walkthrough, techno-managerial, HR fit
Consulting / Big Four	Deloitte, EY, PwC, KPMG, Accenture	Case discussion, client-scenario handling, manager round, values and fit
European employer	Booking.com, Adyen, SAP, Zalando, N26, Revolut	Structured competency rounds, culture and values fit, relocation/visa conversation
GCC / captive centre	Global banks, retail and pharma captives in India	Domain depth, stakeholder management, process and compliance awareness
Regulated / professional	CA firms, tax and audit practices, legal, banking	Technical statute knowledge, ethics and judgement scenarios, articleship fit
Industrial / manufacturing	Automobile OEMs, engineering and supply-chain firms	Core domain fundamentals, plant and process scenarios, safety and quality reasoning

Functional coverage

Functions are grouped into launch waves so that quality is proven before breadth is attempted. Each function requires its own question taxonomy, evaluation rubric and interviewer persona calibration — this is real work per function, not a configuration toggle.

Wave	Functions	Rationale
Wave 1 (launch)	Backend / full-stack engineering; QA and automation; data and AI engineering	Founder has direct domain credibility and an existing audience here, which de-risks question quality and seeds early distribution.
Wave 2 (+3 months)	Product management; management consulting; business and data analysis	High interview-prep spend, strong existing appetite for realistic practice, adjacent enough to validate that the engine generalises beyond engineering.
Wave 3 (+6 months)	Chartered accountancy, taxation and audit; finance and banking	Large, motivated, chronically under-served candidate pool in the Indian market with almost no credible AI prep alternative.
Wave 4 (+9 months)	Automobile and manufacturing; supply chain; HR; sales and marketing	Completes the corporate-function map once the per-function onboarding playbook is repeatable and cheap.

Discipline required here

The ambition is every corporate function. The launch must not be. Each function added before its rubric is properly built dilutes perceived quality across the whole product — and quality is the entire moat. Ship three functions that feel extraordinary rather than twenty that feel generic.

Data strategy: what we source and how

The knowledge layer blends three sources with explicit, user-visible provenance. Provenance is a feature, not a compliance chore — a candidate treats 'reported by 23 candidates in the last 6 months' very differently from 'commonly asked in this type of round', and conflating the two destroys trust.

Tier	Source	Legal posture	Role in product
A Model knowledge	Pre-trained LLM knowledge of interview formats, competency frameworks, domain fundamentals and question archetypes across functions and employers.	Clean. No collection, no storage, no third-party terms engaged.	Baseline coverage for every function and employer from day one. Solves the cold-start problem entirely. Labelled as <i>general pattern</i> .
B Open and permitted	Employer career-site process pages; openly-licensed repositories; official competency frameworks and syllabi (e.g. professional body curricula); individual authors' own blog posts used with attribution or permission.	Defensible where licence or intent permits reuse. Requires a per-source review and a maintained source register.	Grounds employer-specific round structure and official process detail. Labelled as <i>published by employer</i> or <i>attributed source</i> .
C Community contributed	Post-interview submissions from platform users: rounds faced, question themes, difficulty, interviewer behaviour, outcome.	Clean and owned. Contributed under platform terms with explicit licence grant and anonymisation.	The compounding long-term moat. Freshness, corroboration counts and recency weighting. Labelled as <i>reported by N candidates</i> .

Explicitly out of scope

Bulk scraping of LeetCode, Blind, Reddit, Glassdoor, AmbitionBox or GeeksforGeeks discussion content. The reasons are commercial rather than merely cautious: platform terms of service create contract exposure that survives even where scraping public pages is lawful; app-store review and enterprise or campus procurement both ask about data provenance; and a single well-publicised complaint about republished anonymous career content would damage a trust-dependent brand disproportionately. Where a licensed API or partnership is available, that is the route — not a workaround.

The contribution flywheel

Tier C is bootstrapped deliberately rather than hoped for. The mechanism:

- 1 Trigger.** Every candidate who set a target employer receives a prompt shortly after their stated interview date: 'How did it go?'
- 2 Low-friction capture.** A structured two-minute form — rounds faced, themes per round, difficulty rating, interviewer style, outcome. Optionally a voice note that the platform transcribes and structures automatically, which materially raises completion rates over typing.
- 3 Incentive.** Contribution unlocks credits: additional mock interviews, access to a premium report, or early access to a newly-added employer. Contribution is never required to use the product.
- 4 Anonymisation.** Submissions are stripped of identifying detail, aggregated, and surfaced only as patterns with corroboration counts — never as a reproducible individual account.
- 5 Recency decay.** Contributions are weighted by age. A report from three years ago informs the model far less than one from last month, because loops genuinely change.

The candidate profile model

The profile exists for one purpose: to let the engine curate an interview that a real interviewer at that employer, for that role, would plausibly conduct with this specific person. Every field must earn its place by measurably changing question selection or evaluation. Fields that do not are removed — onboarding friction is the primary conversion killer for this product category.

Group	Fields captured	How it changes the interview
Identity and access	Google OAuth identity; display name; preferred interview language (English / Hindi-English code-switched / regional, later).	Determines interviewer language, register and formality level.
Resume (parsed)	Employers and tenure; titles and level progression; projects with technologies; domains; education; certifications; gaps and transitions.	Primary personalisation source. Drives project deep-dive questions, seniority calibration, and probing of gaps or short tenures the way a real panel would.
Function and level	Function (from taxonomy); current level; target level; total experience; people-management scope if any.	Selects question taxonomy and rubric. Sets difficulty band and determines whether leadership or stakeholder rounds are included at all.
Skills	Self-declared skills with confidence rating, reconciled against skills detected in the resume; explicit 'weak areas I want pushed on'.	Balances the interview between strengths and declared weaknesses. Discrepancies between claimed and evidenced skill are themselves probed, as in a real interview.
Session context (<i>per session</i>)	Company and role, named at the start of each session. Optionally round type and interview date.	Resolves employer archetype and composes that session's loop. Deliberately not stored as a persistent target list — see the design note below.
Context and constraints	Location and relocation intent; visa or work-authorisation situation; notice period; compensation expectation band.	Enables realistic rehearsal of relocation, visa, notice and compensation conversations — consistently the least-prepared and highest-anxiety part of European and cross-border loops.
Performance history	Rolling scores per competency; recurring weaknesses; sessions completed; improvement trajectory. Grouped by company and role after the fact.	Makes successive interviews adaptive rather than repetitive. Feeds the practice plan and the derived progress view.

Onboarding sequence. Google sign-in, then resume upload with parsed fields shown for confirmation rather than blank-form entry, then function and level. Everything else is deferred and requested contextually. Company and role are asked at the point of starting an interview, not during onboarding. The target is a first mock interview started within four minutes of landing.

Design note: why there is no stored target list

An earlier draft had the candidate declare up to three concurrent target employers during setup. That was wrong on two counts. It front-loads configuration work before the candidate has received any value, which is the dominant drop-off risk in onboarding; and an arbitrary cap solves a problem the product does not have, since each session is scoped to one company and one role anyway and therefore never mixes requirements. Company and role are chosen per session, with no cap and no setup step. Progress is *derived* from completed sessions grouped by company and role rather than declared in advance — the same retention value with no entity for the candidate to create, edit, or clean up.

Interview engine architecture

A session is a stateful, adaptive conversation rather than a scripted sequence. The engine composes a plan, then departs from it in response to the candidate — which is precisely what makes it feel real.

Session lifecycle

- 1 **Compose.** Resolve profile, the company and role named for this session, archetype and round type into a structured interview plan: opening, competency areas to cover, difficulty ladder, planned depth per area, and time budget.
- 2 **Retrieve.** Pull grounding context from the knowledge layer — employer round structure, function rubric, question archetypes, and any corroborated community reports — with provenance labels attached.
- 3 **Open.** Interviewer persona introduces itself in the candidate's chosen language, sets expectations for the round, and opens with the archetype-appropriate warm-up (a project walkthrough for a service-based loop; a coding framing for a product loop).
- 4 **Conduct.** Turn-by-turn conversation. Every candidate answer is assessed in-flight for depth, correctness and specificity, and the next turn is chosen accordingly: follow up, probe deeper, challenge a weak claim, move on, or raise difficulty.
- 5 **Pressure.** Archetype-appropriate realism — interruption when a candidate rambles, silence after an incomplete answer, pushback on unsupported claims, scenario escalation in case and managerial rounds.
- 6 **Close.** Candidate questions, wrap-up, and a realistic sign-off.
- 7 **Evaluate.** Full transcript scored against the function and level rubric, producing the structured report described in Section 09.

Round types supported

- **Technical fundamentals** — domain concept depth, calibrated to level.
- **Project deep-dive** — interrogation of the candidate's own resume claims. This is the single highest-value round type for service-based and consulting loops and is almost entirely absent from competing products.
- **Coding / practical** — spoken reasoning with an optional editor pane.
- **System or solution design** — product companies and senior levels.
- **Case and client scenario** — consulting, Big Four, service-based delivery.
- **Techno-managerial** — delivery, estimation, escalation, stakeholder handling.
- **Behavioural / competency** — structured, values-based, European-style.
- **HR, fit and closing** — notice, compensation, relocation, visa, motivation.

Voice pipeline: design and unit economics

The interviewer is voice-first. Video avatar is explicitly deferred: real-time lip-synced avatar streaming adds significant latency, cost and engineering risk on top of an already non-trivial pipeline, and delivers a small fraction of the realism gain that voice alone provides. Voice-only is the correct Phase 1 scope.

Pipeline design

The platform composes its own pipeline rather than adopting a bundled conversational-AI product. The orchestration logic — retrieval, adaptive questioning, scoring — is the product itself and cannot be delegated; and self-composing reduces per-minute cost by roughly an order of magnitude.

Stage	Function	Candidate technologies	Selection criteria
Capture	Browser microphone capture, voice-activity detection, barge-in support	WebRTC in-browser; LiveKit or Pipecat for transport orchestration	Sub-second turn detection; must support candidate interrupting the interviewer
STT	Streaming transcription of candidate speech, code-switch capable	Deepgram; Whisper (self-hosted or API); ElevenLabs Scribe	Accuracy on Indian-accented and Hindi-English code-switched speech is the deciding factor, ahead of price
Reasoning	Retrieval, adaptive question selection, in-flight assessment	Frontier LLM via API, with prompt-cached grounding context	Latency and instruction-following on long stateful conversations
TTS	Interviewer voice output in a consistent persona	ElevenLabs Flash (quality benchmark); Cartesia or Deepgram Aura (latency/cost); self-hosted open models (cost floor)	Naturalness under interruption and low time-to-first-audio
Persist	Transcript, turn timings, scores, audio artefacts	Object storage plus relational store	Retention policy and candidate deletion rights

Cost model per one-hour interview

Modelled on published August 2026 rates. Note that in a one-hour interview the interviewer speaks for only a portion of the session — the candidate does most of the talking — so synthesised speech volume is far lower than session duration suggests. This is the key economic asymmetry that makes voice-first viable.

Approach	Basis	Cost / 60-min session	Assessment
Bundled conversational-AI agent	Charged per minute of call duration, roughly \$0.08/min at standard rates and double that beyond concurrency limits	\$4.80 (up to \$9.60 at burst)	Rejected. Pays a large orchestration margin for logic we are building ourselves, and the session clock runs during candidate speech and silence alike.
Self-composed pipeline (recommended)	STT charged per hour of audio; TTS charged per 1,000 characters of interviewer speech only; LLM charged per token	\$1.00 – \$3.00 all-in	Recommended. Interviewer speech in a one-hour session is typically 6,000–9,000 characters, keeping synthesis cost well under a dollar; STT and LLM make up the balance.

Approach	Basis	Cost / 60-min session	Assessment
Self-hosted open models	GPU infrastructure amortised across sessions	Approaching \$0.20 marginal, plus fixed infrastructure and operations	Deferred. Correct at scale; premature before volume justifies the operational burden and quality-tuning effort.

Two cost actions to take before architecting

First, apply to the vendor startup and grant programmes — several major voice providers run programmes offering qualifying early-stage teams substantial free allowances for a year, which could make the entire Phase 1 voice cost negligible. Second, instrument per-session cost from the first day of the private alpha. Voice cost per session is the single number that determines whether the free tier and pricing model in Section 14 survive contact with real usage.

Interviewer persona strategy

The founder's own voice and persona is a genuine differentiator for the Wave 1 engineering functions, where the founder has established audience credibility. It does not transfer automatically. A recognisably engineering-founder persona conducting a chartered accountancy or automobile manufacturing interview undercuts the authority the product depends on.

- **Flagship persona** — founder voice and likeness, engineering functions, Wave 1. Doubles as the marketing hook into the existing audience.
- **Archetype personas** — a small, deliberately-designed set matched to employer archetype and function: a service-based delivery manager, a Big Four engagement manager, a European hiring manager, a professional-practice partner.
- **Consistency requirement** — each persona needs a documented character brief (register, pace, pressure style, follow-up habits) so behaviour stays coherent across sessions. Persona drift is the fastest way to break the realism illusion.
- **Consent and disclosure** — synthetic voice is disclosed to candidates. Any voice cloned from a real person requires that person's explicit written consent and a commercial-use licence from the vendor.

RAG and knowledge layer

Retrieval grounds the interview in employer-specific and function-specific reality. The design priority is provenance-aware retrieval: every retrieved fact carries its tier, recency and corroboration count into the reasoning context, and those attributes surface to the candidate.

Knowledge entities

- **Employer profile** — archetype, typical loop structure, round sequence, evaluation emphases, known process quirks, geography variants.
- **Function rubric** — competency taxonomy per function and level, with observable indicators for each score band.
- **Question archetype** — a reusable question pattern tagged by function, level, round type and competency, with expected follow-up branches and answer indicators.
- **Interview report** — a structured community contribution, anonymised, recency-weighted, aggregated into corroborated patterns.
- **Candidate context** — the parsed resume and performance history, retrieved into every session for personalisation.

Retrieval design decisions

- **Hybrid retrieval** — structured filtering on employer, function, level and round type first, then semantic search within that filtered set. Pure vector search over a mixed corpus retrieves plausible-but-wrong-context questions, which is exactly the failure mode competitors exhibit.
- **Recency weighting** — community-contributed material decays in weight with age. Loops change; stale reports actively mislead.
- **Corroboration thresholds** — a pattern reported once is surfaced as unverified or withheld; patterns corroborated across multiple independent reports are promoted and labelled accordingly.
- **Grounding discipline** — where retrieval returns nothing for a specific employer, the engine falls back to archetype-level and function-level patterns and says so, rather than inventing employer-specific detail. Fabricated specificity is the most damaging possible failure for this product.
- **Prompt caching** — rubrics and employer profiles are stable across sessions and should be cached to control per-session token cost.

Feedback and scoring system

The report is where perceived value concentrates and where conversion to paid happens. A generic encouraging summary destroys credibility; candidates can tell immediately whether feedback is specific to what they actually said.

Report contents

- **Round outcome simulation** — a calibrated judgement on whether this performance would likely advance at this employer and level, with reasoning.
- **Competency scores** — scored against the function and level rubric, each score justified by specific quoted moments from the candidate's own transcript.
- **Answer-level annotations** — the transcript marked up with what worked, what was vague, what a real interviewer would have probed further, and a stronger alternative framing for the weakest answers.
- **Communication analysis** — structure, filler density, pace, rambling, handling of interruption and of not knowing an answer.

- **Targeted practice plan** — the specific competencies to work on before the next attempt, with the next recommended session type.
- **Readiness trajectory** — movement across sessions per target employer, which is the primary retention mechanic and the reason a candidate returns.

Website: information architecture and design

Design direction: modern, restrained, and confidence-inspiring. The emotional context is a nervous candidate preparing for something that matters to their livelihood. The interface should feel calm and serious — closer to a professional tool than a gamified consumer app. Visual noise, playful illustration and aggressive gamification all undercut the realism the product is selling.

Design principles

- **One primary action per screen.** The dashboard's job is to start the next interview. Everything else is secondary.
- **The session screen is nearly empty.** During a live interview the interface shows only what a real interview would: who is speaking, elapsed time, current round, and an exit. No score tickers, no live hints, no distractions.
- **Typography-led, generously spaced.** A single strong typeface family, clear hierarchy, restrained accent colour used for state and action only.
- **Reports are the showcase.** This is where visual richness belongs — charts, annotated transcript, clear competency breakdown.
- **Mobile-responsive from day one,** since a significant share of the Indian audience will arrive on a phone browser well before a native app exists.
- **Accessible by default** — full transcript alongside audio, keyboard navigation, captions, adequate contrast.

Screen inventory

Screen	Purpose	Key elements
Landing	Convert visitor to first session	Value proposition, a real recorded sample interview clip, coverage breadth (service-based and non-tech shown prominently as the differentiator), single sign-in call to action
Onboarding	Profile to first interview in under four minutes	Google sign-in, resume upload with parsed-field confirmation, function and level, first target employer, language preference
Dashboard	Orient and start the next session	Single prominent start-interview action; recent sessions grouped by company and role with trend; recent reports; contribution prompt
Session setup	Name the interview context, in seconds	Company search with archetype resolution, role, optional round type; expected loop preview. Reached from the start-interview action, not from onboarding
Pre-session	Set expectations and check hardware	Round type, duration, language, microphone test, persona introduction
Live session	Conduct the interview	Minimal interface: speaking indicator, timer, round label, optional live transcript toggle, optional code editor pane, end-session control
Report	Deliver the value and drive return	Outcome simulation, competency scores, annotated transcript, practice plan, next-session call to action
Progress	Retention through visible improvement	Competency trend lines; session history grouped by company and role, derived from completed sessions rather than declared targets
Contribute	Feed the flywheel	Two-minute structured form or voice note, credit reward, anonymity assurance

Screen	Purpose	Key elements
Profile	Maintain personalisation inputs	Resume re-upload, skills reconciliation, constraints, language, account and data-deletion controls

11 Technology stack

Stack selection deliberately favours the founder's existing operating experience over theoretically optimal choices. A solo or very small team shipping a real-time voice product cannot also absorb the learning curve of an unfamiliar backend ecosystem.

Layer	Recommendation	Rationale
Frontend	Next.js with TypeScript; Tailwind; a headless component library	Server rendering for landing-page performance and SEO, which matters because organic search on employer-plus-interview queries is a primary acquisition channel
Backend	Spring Boot (Java/Kotlin)	Directly matches founder expertise, minimising delivery risk on the most complex part of the system. Strong fit for stateful session orchestration
Real-time transport	LiveKit or Pipecat alongside WebRTC	Purpose-built for low-latency voice agents; avoids rebuilding turn-taking, barge-in and transport handling from scratch
Data store	PostgreSQL with the pgvector extension	Single store for relational and vector workloads at this scale. Avoids operating a separate vector database prematurely
Managed platform	Supabase for auth, storage and Postgres	Proven in the founder's prior product; Google OAuth, storage and row-level security available immediately
Async processing	Queue-backed workers for transcript scoring and report generation	Report generation is expensive and must not block session teardown
Observability	Per-session cost, latency and quality instrumentation from day one	Unit economics and turn latency are the two numbers that determine viability
Payments	Razorpay for India; Stripe for international	Matches the target market's payment behaviour, including UPI

12 Phased delivery roadmap

Phases are gated on evidence, not on elapsed time. Each gate is a specific question that must be answered affirmatively before the next phase begins.

Phase	Scope	Exit gate
Phase 0 Private alpha ~6 weeks	Voice pipeline proven end to end. One function (backend engineering), two employer archetypes (one global product, one service-based). Resume parsing, single target, one round type, basic report. Manually seeded knowledge base.	Do twenty hand-picked testers describe the session as feeling like a real interview? If not, nothing else matters yet.
Phase 1 Public web launch ~10 weeks	Wave 1 functions. All eight employer archetypes. Per-session company and role selection. Full round-type coverage. Complete report and progress tracking. Hindi-English mode. Payments and free tier live.	Positive unit economics per paid session, and week-four retention indicating candidates return for a second and third session.
Phase 2 Flywheel and breadth ~12 weeks	Contribution flow live with credits and anonymisation. Corroboration and recency weighting in retrieval. Wave 2 functions. Additional archetype personas. SEO content engine on employer-plus-round queries.	Meaningful contribution rate from completed interviews, and community-grounded sessions rating measurably higher than baseline sessions.

Phase	Scope	Exit gate
Phase 3 Mobile <i>~10 weeks</i>	Android application first, iOS second. Full session capability on mobile, push-based practice reminders, offline report reading.	Sustained mobile web traffic share and explicit user demand confirming an app is the constraint rather than a vanity milestone.
Phase 4 Expansion <i>ongoing</i>	Waves 3 and 4 functions. Video avatar layer. Institutional offering for colleges and training providers. Regional language support.	Repeatable per-function onboarding playbook, and a validated institutional sales motion.

Web-to-mobile transition strategy

Because a native application is a stated Phase 3 objective, the web build should be structured now to make that transition cheap. Retrofitting is expensive; the decisions below cost almost nothing if taken at the start.

- **Build against a clean public API from day one.** The web client consumes exactly the same versioned API the mobile clients will. No server-rendered coupling in session, report or profile flows.
- **Keep all interview logic server-side.** Composition, retrieval, scoring and persona behaviour live in the backend, never the client. The mobile app should be a thin client over the same engine.
- **Choose a voice transport with mature mobile SDKs.** This is the single most consequential Phase 3 decision made during Phase 1 — a transport layer without solid Android and iOS support forces a rewrite of the hardest component.
- **Ship the web app as an installable PWA.** This delivers a home-screen installable experience, captures mobile demand during Phases 1 and 2, and generates the usage evidence that tells you whether the native app is actually warranted.
- **Design mobile-first for the session screen specifically.** The live interview interface is minimal by design and therefore translates almost unchanged to mobile, which makes it the natural first surface to port.
- **Front-load store compliance thinking.** Microphone permission rationale, subscription handling through the store billing systems, and account-deletion requirements all shape API design and should be anticipated rather than discovered during review.
- **Android before iOS.** Matches the target market's device distribution and carries materially lower process overhead, allowing iOS to follow once the mobile experience is validated.

Monetisation and unit economics

Pricing must reflect a real marginal cost per session, which distinguishes this product from text-based competitors that can afford generous free tiers. The free tier exists to demonstrate quality once, not to provide ongoing free practice.

Tier	Offer	Intent
Free	One full mock interview per month, complete report, any company and role	Demonstrate realism. The report is deliberately not truncated — conversion depends on the candidate seeing full value once
Focused	Four sessions per month, progress tracking, full round-type access, any company and role	The core individual plan, aimed at a candidate in an active interview process
Intensive	Unlimited fair-use sessions, priority scheduling, all personas, detailed benchmarking	For candidates in a compressed high-stakes loop, where willingness to pay is at its peak
Credits	Pay-as-you-go session packs	Serves occasional users and provides the reward currency for community contributions
Institutional	Seat-based licensing for colleges, training institutes and placement cells, with cohort analytics	Higher contract value, predictable revenue, and a natural fit with the founder's existing teaching and mentoring channels

Pricing posture. Price in local currency with India-first anchoring and regional purchasing-power adjustment, consistent with the founder's prior product experience. Verify gross margin per tier against measured per-session cost before publishing prices, not after.

Success metrics

Category	Metric	Why it matters
Realism (north star)	Share of completed sessions rated as feeling like a real interview	This is the moat expressed as a number. If it falls, nothing downstream recovers
Activation	Sign-up to first completed session, and time to first session	Onboarding friction is the dominant early drop-off
Completion	Share of started sessions completed in full	Abandonment mid-interview signals latency problems, persona failure or mis-calibrated difficulty
Retention	Candidates completing a second and third session within four weeks	Proves the report drives return, which is the entire retention thesis
Efficacy	Self-reported interview outcomes among candidates who practised	The strongest marketing asset available, and the hardest to fake
Flywheel	Contribution rate per completed interview; corroborated patterns added	Measures whether the long-term data moat is actually compounding
Economics	Fully-loaded cost per session; gross margin per paid tier	Determines whether the free tier and pricing survive real usage
Latency	Median turn-around time from candidate finishing speaking to interviewer responding	Above roughly a second, realism collapses regardless of content quality
Breadth	Session quality parity across functions and archetypes	Guards against non-engineering coverage silently becoming second-class

Risks, legal posture and mitigations

Risk	Assessment	Mitigation
Data provenance and platform terms	High impact. Scraping career-content platforms creates contract exposure independent of whether scraping public pages is lawful, and threatens store review and institutional sales.	Tiered sourcing per Section 04. Maintain a source register. Pursue licensed partnerships rather than workarounds. Obtain a legal review of the source register before public launch.
Realism falls short	Existential. The entire differentiation rests on the session feeling real.	Gate Phase 0 exit on the realism rating. Invest in conversation design and latency ahead of breadth. Maintain documented persona briefs to prevent drift.
Fabricated employer specifics	Severe credibility damage. A confidently invented claim about a named employer's process is worse than admitting ignorance.	Strict grounding discipline with archetype-level fallback and explicit provenance labelling. Never present a general pattern as a specific report.
Voice cost overrun	Moderate. Marginal cost per session directly constrains the free tier.	Self-composed pipeline. Per-session cost instrumentation from day one. Apply to vendor startup grant programmes. Keep a self-hosted fallback path open.
Cold start on breadth	Moderate. Non-engineering functions have no seeded audience.	Wave-based rollout. Do not launch a function until its rubric and persona are built. Use the founder's existing channels to seed Wave 1 before expanding.

Risk	Assessment	Mitigation
Candidate data sensitivity	High. Resumes, interview recordings and outcomes are unambiguously personal data.	Explicit consent for recording. Clear retention limits. Self-service deletion. Anonymisation before any contribution enters the shared corpus. Regional data protection obligations assessed before European users are onboarded.
Incumbent response	Moderate. A well-funded competitor could add voice.	Compete on the axes that are expensive to copy — language, employer-archetype breadth and interview realism — rather than on feature checklists.

Open decisions

These require a decision before or during Phase 0. Each materially changes the build.

#	Decision	Considerations
1	Founder voice clone as the flagship persona, or a designed persona in the founder's style?	A clone maximises the audience hook and authenticity, but ties brand identity to one function's credibility and creates a personal-likeness dependency. A designed persona scales more cleanly across functions.
2	Which two employer archetypes anchor the Phase 0 alpha?	The recommendation is one global product and one service-based, because proving the engine handles both extremes validates the architecture. Choosing two similar archetypes would prove far less.
3	Is Hindi-English code-switched mode in Phase 0 or Phase 1?	It is the sharpest differentiator and the most audience-resonant feature, but it raises the STT accuracy bar considerably. Recommendation is to prototype in Phase 0 and ship in Phase 1.
4	Does the free tier include a full report?	A truncated report lifts short-term conversion but undermines the quality demonstration that the whole acquisition strategy depends on. Recommendation is full report, strictly limited frequency.
5	Institutional offering in Phase 2 or Phase 4?	The founder has existing college and mentoring channels that could produce early revenue and a captive contribution base, at the cost of pulling focus from the consumer product during its most fragile period.
6	Build a code-execution environment, or spoken reasoning only?	Spoken reasoning matches how many real interviews of this type actually run and is far cheaper to build. A real editor is table stakes for product-company coding rounds specifically.
7	What is the recording retention policy?	Longer retention improves the product through better scoring calibration; shorter retention reduces risk and is easier to justify to candidates. This decision affects onboarding consent copy and must be settled before the alpha.

Recommended immediate next step

Do not begin the full build. Spend the first two weeks proving the single riskiest assumption in this document: that a self-composed voice pipeline can conduct a sub-second-latency, code-switched, adaptive interview that a real candidate rates as realistic. Build that one vertical slice — one function, one employer, one round type, no payments, no dashboard, no polish. Every other section of this PRD is conventional product engineering that is known to be achievable. That slice is not, and it determines whether the product is worth building at all.

This document reflects publicly available pricing and platform information as of August 2026 and represents product and engineering judgement, not legal advice. The data-sourcing posture in Section 04 and the personal-data handling in Section 16 should be reviewed by qualified counsel before public launch.